RF Prior: Context-Preserving Representation Adaptation for Instance Segmentation Transfer

Jason K. Nam[1,2]
[1] Department of Data Science, Hongik University
[2] AI Research Division, ETRI
skarlqja68@gmail.com

Abstract

Instance segmentation transfer requires learning target boundaries while retaining context encoded by a pretrained model. We introduce RF PRIOR*, a receptive-field (RF)-constrained framework that couples scale-selective preservation with a trainable contextual readout. A scale-selective parameter constraint preserves source-trained P5 weights as upstream P3/P4 features adapt along its gradient path. Pooled P5 context guides P4 through area-restricted attention (AA), followed by local P3 refinement. On three sidewalk targets, this design improves mask mean average precision (mAP) over scratch and full fine-tuning at 25.2 frames per second (FPS). Across three targets and seeds, attention reduces P4 displacement and late gradients as P3 adapts and P5 features evolve. Shared-head losses, including distribution focal loss (DFL), connect this scale-specific adaptation to mask and localization fitting. These results connect contextual reuse to preserved parameters and their target-adaptive readout. The project page is available at* https://namkibeom.github.io/rf-prior/.

1. Introduction

Transferring pretrained features to instance segmentation changes the spatial evidence required for prediction. Target masks demand both contextual recognition and precise boundaries, while changes in label inventory and annotation geometry reshape the supervision across scales [27, 53]. Because a feature pyramid distributes these cues across receptive fields [32], transfer poses a coupled question: which source parameters should be preserved, and how should their features interact with target-adaptive spatial detail?

Full fine-tuning, starting-point ℓ_2 regularization (L2-SP), and selective tuning regulate departure from source parameters [21, 23]. Yet parameter control leaves the interaction between evolving features and downstream fusion unspecified. Frozen-backbone detection establishes readout capacity as a key factor [43]. Dense transfer therefore requires coupling preservation to a decoder that integrates retained context with target boundaries and foreground cues.

RF PRIOR couples preserved P5 parameters to trainable P3/P4 in a You Only Look Once (YOLO)v12-X pyramid [41]. As P5 inputs evolve, gradients traverse the preserved transformation and couple it to target learning. Spatial pyramid pooling-fast (SPPF) exposes contextual scales for AA to integrate with P4, after which P3 refines local boundaries and foreground geometry.

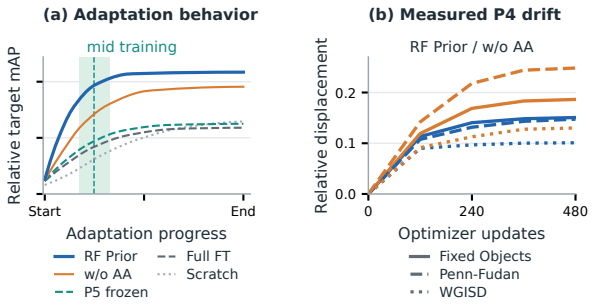


Figure 1. **Adaptation behavior and measured dynamics.** (a) Early-adaptation illustration; mid training is an intermediate target-adaptation stage. (b) Recorded P4 displacement, averaged over three seeds: AA reduces final movement by 19.2%/40.6%/22.4% versus w/o AA. The reduction holds across all nine matched target–seed pairs at the 480-update endpoint.

We evaluate this design on SideGuide's POLYSEG/SURFACEMASK targets [34] and FIXED OBJECTS, which expose changes in boundary supervision and label inventory. Three-target diagnostics connect these results to parameter movement and feature evolution. Our contribution is a preservation–readout formulation that separates fixed parameters from evolving representations, together with stage-resolved evidence linking this distinction to target-mask learning across spatial scales.

2. Related Work

Parameter preservation and feature adaptation. L2-SP anchors source weights; Learning without Forgetting (LwF) and elastic weight consolidation (EWC) constrain source behavior or important parameters [19, 23, 25]. SpotTune selects fine-tuning paths per input, and surgical fine-tuning relates layer choice to shift type [7, 21]. Residual

adapters, visual prompts, and Rein adapt features around preserved backbones [13, 37, 51]; weight-decomposed low-rank adaptation (DoRA) separates magnitude and direction [28]. These works establish that fixed weights support changing representations. We specialize this principle to a spatial hierarchy: preserved wide-context transformations receive adaptive inputs, and the P5-to-P4 readout shapes local-stage gradients. Stage-resolved diagnostics distinguish this coupling from weight anchoring.

Dense transfer across tasks and scenes. Feature reuse depends on task distance, depth, and spatial support [32, 53]. Where frozen-backbone detection emphasizes readout capacity [43], our lower pyramid also adapts to target geometry. SideGuide's box and polygon annotations expose supervision and inventory shifts within related sidewalk scenes [34], complementing street-scene benchmarks [4, 33, 54] and synthetic-to-real or geographic adaptation [9, 11, 44]. Penn-Fudan pedestrians and the Wine Grape Instance Segmentation Dataset (WGISD) extend these diagnostics across target geometries [39, 48].

Contextual readout for dense prediction. Fully convolutional and pyramid models combine spatial scales [27, 30, 56]; Mask region-based convolutional neural networks (Mask R-CNN), You Only Look At CoefficienTs (YOLACT), and Segmenting Objects by Locations v2 (SOLOv2) realize different mask readouts [1, 8, 50]. Non-local, criss-cross, and transformer operators expand spatial interaction [2, 10, 49, 58]. SegFormer, Mask2Former, MaskDINO, and OneFormer develop hierarchical or masked reasoning [3, 12, 22, 29, 52]; YOLO variants emphasize real-time prediction [38, 45–47]. OneFormer with a large shifted-window Transformer backbone (OneFormer-Swin-L) supplies our accuracy reference. Our P5-to-P4 interface couples efficient cross-scale processing [57] to preserved context, followed by lightweight P3 refinement of target mask boundaries.

Automatic box-to-polygon priors. Box-supervised segmentation derives mask supervision through proposals, structured teachers, or consistency losses [5, 17, 20, 42]. Our offline conversion uses Segment Anything Model (SAM) proposals [18], intersection-over-union (IoU) gating, clipping, and contour simplification [6, 36]. Every FIXED OBJECTS variant uses these polygons, holding supervision fixed across parameter and readout designs.

3. Methodology

Preliminaries. Given $x \in \mathbb{R}^{3\times H\times W}$, the backbone $B_{\boldsymbol{\theta}_{\mathrm{bb}}}$ produces a feature pyramid $\{F_3, F_4, F_5\}$ at strides $\{8, 16, 32\}$. Blocks $\boldsymbol{\theta}_{\mathrm{bb}} = [\boldsymbol{\theta}_3; \boldsymbol{\theta}_4; \boldsymbol{\theta}_5]$ end at P3/P4/P5 with nominal receptive-field scales $R_3 < R_4 < R_5$. The head $H_{\boldsymbol{\theta}_{\mathrm{seg}}}$ fuses these features into instance masks. Detector parameters $\boldsymbol{\theta}_0$ are learned on bounding-box source $\mathcal{D}_{\mathrm{bb}}$ (*BBox-DS*); $\mathcal{D}_{\mathrm{seg}} \in \{\text{POLYSEG}, \text{SURFACEMASK}, \text{FIXED OBJECTS}\}$ supplies polygons for target instance segmentation.

Definition 1 (Scale-selective preservation)**.** Let $\boldsymbol{\Theta} = [\boldsymbol{\theta}_3; \boldsymbol{\theta}_4; \boldsymbol{\theta}_5]$ and $\mathcal{C} = \{\boldsymbol{\Theta} : \boldsymbol{\theta}_5 = \boldsymbol{\theta}_{0,5}\}$. On $\mathcal{C}$, P3/P4 and the readout remain trainable; P5 inputs and normalization state evolve under preserved learned parameters.

Cross-context objective. Set $\boldsymbol{m} = [\mathbf{0}; \mathbf{0}; \boldsymbol{\theta}_{0,5}]$ and block precision $\boldsymbol{\Lambda}_\tau = \mathrm{blkdiag}(\xi_3 I, \xi_4 I, \tau I)$. For $\xi_3, \xi_4, \tau > 0$, the Gaussian family

$$p_\tau(\boldsymbol{\Theta}) \propto \exp\left[-\tfrac{1}{2}(\boldsymbol{\Theta} - \boldsymbol{m})^\top \boldsymbol{\Lambda}_\tau (\boldsymbol{\Theta} - \boldsymbol{m})\right], \tag{1}$$

converges as $\tau \to \infty$ to P3/P4 Gaussian shrinkage and $\delta(\boldsymbol{\theta}_5 - \boldsymbol{\theta}_{0,5})$. For targets $\mathcal{D}_{\mathrm{seg}} = \{(x_i, y_i)\}_{i=1}^N$ and task loss in Eq. (9), the cross-context penalized risk is

$$\begin{aligned}\mathcal{L}_{\mathrm{CC}}(\boldsymbol{\theta}) = \sum_i \ell_{\mathrm{task}}\big(H_{\boldsymbol{\theta}_{\mathrm{seg}}} \circ B_{\boldsymbol{\theta}_{\mathrm{bb}}}(x_i), y_i\big) \\ + \frac{\xi_3}{2}\|\boldsymbol{\theta}_3\|_2^2 + \frac{\xi_4}{2}\|\boldsymbol{\theta}_4\|_2^2, \quad \boldsymbol{\theta}_5 = \boldsymbol{\theta}_{0,5}.\end{aligned} \tag{2}$$

L2-SP instead anchors the source vector $\boldsymbol{\Theta}_0$ through

$$\Omega_{\mathrm{SP}} = \tfrac{\lambda}{2}\|\boldsymbol{\Theta} - \boldsymbol{\Theta}_0\|_2^2, \quad \nabla\Omega_{\mathrm{SP}} = \lambda(\boldsymbol{\Theta} - \boldsymbol{\Theta}_0), \tag{3}$$

which restores source proximity while allowing every block to adapt [23]. For stochastic gradient descent from $\boldsymbol{\Theta}^t \in \mathcal{C}$, set $g^t = \nabla_{\boldsymbol{\Theta}}\ell^t + \boldsymbol{\Lambda}_0\boldsymbol{\Theta}^t$ and $\boldsymbol{\Lambda}_0 = \mathrm{blkdiag}(\xi_3 I, \xi_4 I, 0)$. Definition 1 gives

$$\begin{aligned}\boldsymbol{\Theta}^{t+1} &= \Pi_{\mathcal{C}}(\boldsymbol{\Theta}^t - \eta_t g^t) \\ &= \boldsymbol{\Theta}^t - \eta_t P_{\mathcal{T}} g^t.\end{aligned} \tag{4}$$

Here $\Pi_{\mathcal{C}}(u) = \boldsymbol{m} + P_{\mathcal{T}} u$, with $P_{\mathcal{T}} = \mathrm{blkdiag}(I, I, 0)$, removes the P5 update direction. The same constraint removes the P5 component of each AdamW displacement in our matched 480-update transfer studies [31].

Global-to-local gradient transmission. Write $f_\ell = \mathrm{vec}(F_\ell)$, $J_{ab} = \partial f_a / \partial f_b$, and $J_{\ell\theta} = \partial f_\ell / \partial \boldsymbol{\theta}_\ell$. Let g_ℓ be the loss gradient at each separate readout input and $h_4 = g_4 + J_{54}^\top g_5$.

Proposition 1 (Frozen-stage gradient transmission)**.** *At a differentiable point, retaining P5 in the computation graph under $\mathcal{C}$ yields*

$$\begin{aligned}\nabla_{\boldsymbol{\theta}_4}\ell &= J_{4\theta}^\top\big(g_4 + J_{54}^\top g_5\big), \\ \nabla_{\boldsymbol{\theta}_3}\ell &= J_{3\theta}^\top\big(g_3 + J_{43}^\top h_4\big).\end{aligned} \tag{5}$$

The contextual gradient term $J_{4\theta}^\top J_{54}^\top g_5$ is retained even though the P5 parameter update is zero.

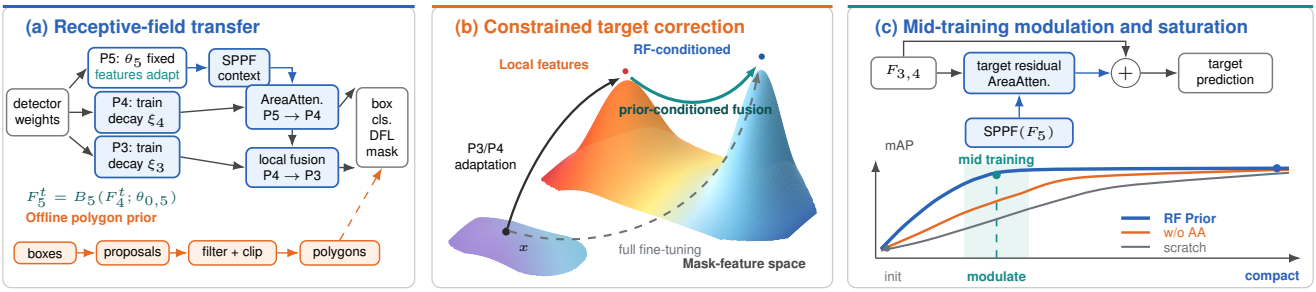


Figure 2. **Overview of RF PRIOR.** (a) P5 preserves source-trained parameters while its feature maps evolve with the trainable input representation. P3/P4 receive weight decay, and gradients pass through P5. (b) Conceptual mask-feature landscape: SPPF/AA provides prior-conditioned fusion; surface height is illustrative. (c) Schematic mid-training modulation and saturation: adapted context can facilitate subsequent refinement. Figure 4 offers a qualitative view of retained source-conditioned structure after FIXED OBJECTS adaptation.

The readout thus shapes both stages through distinct Jacobians, derived in the supplement.

Vectorize spatial features as $X_\ell \in \mathbb{R}^{n_\ell \times c_\ell}$. Resampling $S_{5\to\ell} \in \mathbb{R}^{n_\ell \times n_5}$ and channel maps $\Phi = \{W_{5\ell}, W_\ell\}_{\ell\in\{3,4\}}$ give a common d-channel basis [27]:

$$\begin{aligned} R_{\ell,i} &= S_{5\to\ell} X_{5,i} W_{5\ell} - X_{\ell,i} W_\ell, \\ E_{\text{align}}(\Phi) &= \sum_{i,\ell} \|R_{\ell,i}\|_F^2 = \sum_{i,\ell} \operatorname{tr}(R_{\ell,i}^\top R_{\ell,i}). \end{aligned} \tag{6}$$

For $\ell \in \{3,4\}$, the trace separates feature energies from cross-scale association at specified projections; target supervision learns the contextual fusion used for prediction.

Multi-scale attentive decoder. Let $P_0 = X_5 W_c$ and apply stride-one 5×5 max pooling $\mathcal{M}_5$ on the underlying spatial grid, returning vectorized features:

$$\begin{aligned} P_j &= \mathcal{M}_5^j(X_5 W_c), \quad k_j = 4j+1, \\ \widetilde{X}_5 &= [P_0 \| P_1 \| P_2 \| P_3] W_s \in \mathbb{R}^{n_5 \times d_s}. \end{aligned} \tag{7}$$

For $j \in \{1,2,3\}$, k_j gives the 5/9/13 interior pooling widths; $W_s \in \mathbb{R}^{4d_c \times d_s}$ fuses these contexts [14]. We upsample $\widetilde{X}_5$ and form the stride-16 matrix $U_4 = [X_4 W_4 \,\|\, S_{5\to4}\widetilde{X}_5]$. Following the area-restricted interaction used by YOLOv12 [41], for region index set $\mathcal{I}_r$ attention computes

$$\begin{aligned} [Q_r \,\|\, K_r \,\|\, V_r] &= U_4[\mathcal{I}_r][W_Q \,\|\, W_K \,\|\, W_V], \\ A_r &= \operatorname{softmax}(Q_r K_r^\top/\sqrt{d}), \quad Z_r = A_r V_r. \end{aligned} \tag{8}$$

Reassembled outputs $\{Z_r W_O\}_{r=1}^a$ enter a P4 residual projection. Context shapes both affinities A_r and values V_r, hence Proposition 1's gradient; a two-convolution cross-stage-partial (C2f)-style block then connects P4 to P3. For n stride-16 tokens, a equal regions reduce attention cost from $O(n^2 d)$ to $O(n^2 d/a)$. The w/o-AA ablation retains P5 preservation and SPPF with local fusion; plain "P5 Frozen" retains the original neck and its feature-fusion operators.

Unified optimization. The YOLO-style multitask loss combines complete intersection over union (CIoU), binary cross-entropy (BCE), distribution focal loss (DFL), and mask supervision [1, 24, 59]:

$$\mathcal{L} = \lambda_{\text{box}}\mathcal{L}_{\text{CIoU}} + \lambda_{\text{cls}}\mathcal{L}_{\text{BCE}} + \lambda_{\text{dfl}}\mathcal{L}_{\text{DFL}} + \lambda_{\text{mask}}\mathcal{L}_{\text{mask}}. \tag{9}$$

With area-normalized BCE over matched instance pixels, mask supervision contributes alongside the other heads at shared features, coupling localization, class separation, and discrete offset distributions to the contextual readout.

Polygon-prior mining. For proposal M_k and source box B, define $q_k = \text{IoU}(M_k, B)$ and the accepted set $\mathcal{G}_\tau = \{k : q_k \geq \tau\}$. The path used to create FIXED OBJECTS is

$$M^\star = \begin{cases} \bigvee_{k\in\mathcal{G}_\tau} M_k, & \mathcal{G}_\tau \neq \varnothing, \\ M_{\arg\max_k q_k}, & \text{otherwise}, \end{cases} \tag{10}$$

SAM supplies box-prompted candidates [18]. We simplify the union's largest external contour, clamp vertices to B, and use the rectangle for an empty contour. Mining precedes optimization: the retained labels are held fixed and shared by every variant. Table 1(b) reports original-label visual consistency across the retained validation archive and proposal-mask consistency on 100 images.

4. Experiments

Datasets and setup. *BBox-DS* supplies the detector prior; POLYSEG changes box supervision to polygons, SURFACEMASK introduces surface-centric labels, and the smaller FIXED OBJECTS set emphasizes stationary infrastructure (Table 1). Main YOLOv12-X runs use `optimizer=auto` and mixed precision on NVIDIA RTX A6000 (POLYSEG, SURFACEMASK) or A100 (40 GB; FIXED OBJECTS) GPUs. Table 2 FPS is measured on the A100. Evaluation reports precision, recall, and mean average precision (mAP) for boxes and masks; $\text{mAP}_{50:95}$ averages IoU from 0.50 to 0.95 in 0.05 increments. Matched

variants share target data, augmentations, evaluation, and schedule; only source initialization, parameter constraint or regularization, and decoder fusion change. The standard SideGuide comparison uses 30 epochs and the extended row uses 50 epochs. For FIXED OBJECTS, the reported budgets are 16.98k and 50.94k mini-batch steps, with batch size 8 and 1,698 steps per epoch. Subscripts mark second-best recorded checkpoints; budget labels follow the archived convention. Figure 5(d) separately pairs the FIXED OBJECTS logs at identical recorded epochs.

Definition 2 (Stage-wise adaptation profile)**.** For nonzero initial weights, define displacement $\Delta_\ell^\theta(t) = \|\boldsymbol{\theta}_\ell^t - \boldsymbol{\theta}_\ell^0\|_2/\|\boldsymbol{\theta}_\ell^0\|_2$. Pre-clipping gradient root mean square (RMS) is $G_\ell(t) = \|\nabla_{\boldsymbol{\theta}_\ell}\ell^t\|_2/\sqrt{d_\ell}$ for $d_\ell > 0$ recorded elements. Feature distance $D_\ell^{\text{feat}}(t)$ averages $1 - \cos(\bar{F}_{\ell,i}^t, \bar{F}_{\ell,i}^0)$ over fixed images after spatial pooling and vectorization of each feature tensor.

These quantities describe complementary aspects of the same trajectory. Displacement accumulates departure from initialization, whereas gradient RMS records the local training signal before clipping. Feature RMS measures activation magnitude, and cosine distance measures directional change on the same image set. Consequently, a preserved parameter block can have zero displacement while its feature statistics evolve as its inputs change. Comparing the three stages under matched updates locates where each transfer policy absorbs target supervision within the adapting spatial hierarchy of the transferred backbone.

Testing preservation and readout. Scratch, full fine-tuning, plain P5 freezing, and L2-SP test initialization and Definition 1's parameter constraint; full/w/o-AA decoders test the readout under fixed P5. Definition 2 links these controls to stage-wise dynamics, while YOLO/OneFormer references and shared MobileSAM candidates [55] support accuracy–throughput and label-consistency comparisons.

Transfer accuracy and adaptation budget. Table 2 evaluates the coupled preservation–readout design against L2-SP [23] and established segmentation references [12]. On FIXED OBJECTS, the full decoder exceeds the strongest non-proposed result in each column by 1.87 box mAP_{50}, 2.69 mask mAP_{50}, 1.84 box $\text{mAP}_{50:95}$, and 2.19 mask $\text{mAP}_{50:95}$. RF PRIOR is strongest on POLYSEG, improving mask mAP_{50} by 3.56 points over L2-SP and 4.17 points over full fine-tuning. On SURFACEMASK, YOLOv9-E leads coarse mAP_{50}, whereas RF PRIOR leads strict $\text{mAP}_{50:95}$, indicating a strict-localization advantage for already plausible surfaces. The OneFormer-Swin-L reference is competitive but slower: RF PRIOR is 3.4× faster (25.2 vs. 7.4 FPS). Against No Pretraining, it raises compute from 162 to 393 GFLOPs (2.43×) while FPS decreases by 2.3%; SRI 2.34 reflects realized arithmetic throughput, and GFLOPs expose the higher arithmetic cost. The gains are largest for POLYSEG boundaries and FIXED OBJECTS mask prediction under limited adaptation budgets.

On SideGuide, the 30-epoch comparison in Table 3 gives RF PRIOR the best recall and box/mask mAP among matched YOLOv12-X variants. Plain P5 freezing improves strict localization over full fine-tuning but remains below scratch; the complete decoder closes that gap, showing the importance of reading out the preserved context. The 50-epoch row improves every box and mask endpoint under the extended schedule on the same held-out validation split.

Contextual readout across training budgets. SPPF exposes several context radii; AA weights their fusion with P4, whereas RF PRIOR (w/o AA) mixes them locally [14, 41]. At the archived early checkpoint in Table 4, the full decoder improves over scratch by +3.62/+4.44 box/mask mAP_{50} and +2.67/+3.26 box/mask $\text{mAP}_{50:95}$. The lead over RF PRIOR (w/o AA) is 2.18/1.31 box/mask mAP_{50} and 1.64/0.73 strict mAP. At the extended checkpoint, w/o AA leads mask mAP_{50}/$\text{mAP}_{50:95}$ by 0.46/0.18 points, while full retains a 0.29-point strict-box edge. To examine this budget dependence, paired-epoch diagnostics track optimization under the same supervision.

Qualitative protocols. Figure 3 compares direct-target P3–P5 feature maps; coral guides mark matched annotation coordinates. In contrast, Figure 4 follows *BBox-DS* pretraining and FIXED OBJECTS adaptation; scratch uses FIXED OBJECTS alone. Its cross-target POLYSEG/SURFACEMASK rows examine retained source-conditioned structure after the label-inventory change, while the self-target FIXED OBJECTS row reveals spatial support for labeled infrastructure after target adaptation.

Shared-head optimization and accuracy. Figure 5(d) pairs validation mask-loss reduction and mask $\text{mAP}_{50:95}$ difference at 33 epochs: Pearson $r = 0.694$ overall and $r = 0.635$ after warm-up. The full decoder lowers mask loss in 25 epochs and improves strict mask mAP in nine, distinguishing pixelwise fitting from instance ranking.

Distributional localization provides a complementary view of the same adapted features. DFL learns discrete box-offset distributions [24]; mask and box heads share the fused pyramid, so contextual refinement can affect both foreground support and localization. Where D-FINE explicitly refines distributions [35], our unchanged DFL objective measures representation-induced effects. SideGuide improves all four loss families in both phases (Figure 5(a)). Validation DFL retains $0.10815/0.39044 = 27.7\%$ of its early advantage, versus 11.7% for mask loss. Normalized train/validation geometric-mean reductions are

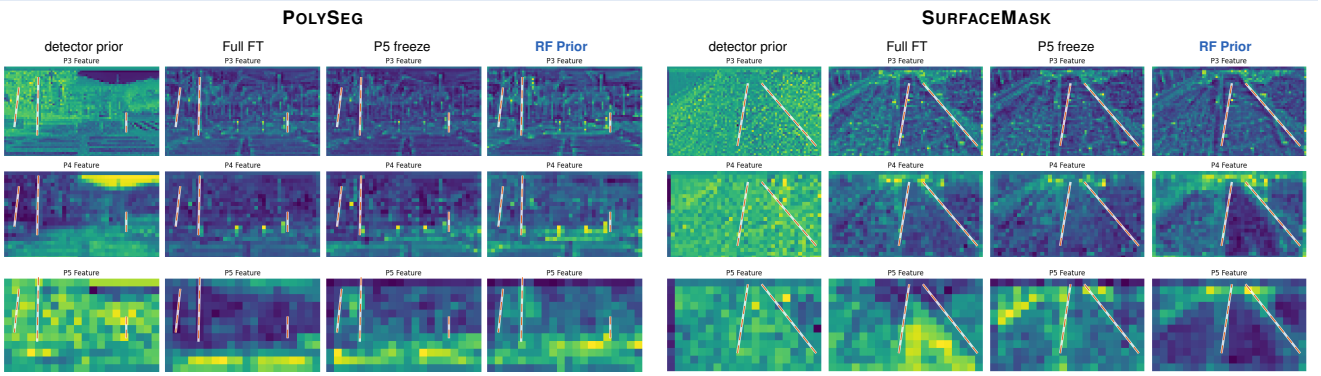


Figure 3. **Direct-target P3–P5 feature maps.** The detector prior precedes adaptation; full fine-tuning (Full FT), P5 freeze, and RF PRIOR use the pictured target. Coral guides reproject Figure 4 input geometry: tree support, signal mast, and curb bollard for POLYSEG; green-strip boundaries for SURFACEMASK, at corresponding spatial positions across the three pyramid levels for matched spatial comparison.

Table 1. **Datasets and FIXED OBJECTS label audits.** (a) BG/LS: background/label-space shift from *BBox-DS*; ▲ mild, ✓ clear. (b) MobileSAM proxy-mask overlap on 100 screened images (IoU: intersection over union; boundary F1: boundary precision–recall harmonic mean; DP: Douglas–Peucker). Final row: AI visual consistency / unresolved / flagged rates (%) for 1,599 retained images.

(a) Dataset composition

Datasets (#images)	Dataset Type	classes	*train*	*validation*	*total*
SIDEGUIDE [34]					
BBox-DS	Prior source	29	20,097	2,233	22,330
POLYSEG	BG ▲, LS ▲	29	66,082	7,342	73,424
SURFACEMASK	BG ✓, LS ✓	6	41,759	4,640	46,399
FIXED OBJECTS (Proposed)	BG ▲, LS ✓	15	13,577	1,543	15,120

(b) FIXED OBJECTS label audits

Selection / post-processing	Mask IoU (%)	Boundary F1 (%)
SAM confidence	85.18±15.99	72.76±25.20
IoU-selected mask	88.41±12.61	76.68±22.62
IoU + coverage	87.63±12.62	73.50±24.35
IoU-gated union	88.75±12.20	77.23±21.53
Union + clip + DP	**88.87±13.05**	**80.14±19.00**
Visual consistency (Astra)	96.1 / 3.1 / 0.8%	

Table 2. **Transfer and efficiency comparison.** Gray marks proposed configurations; bold/underline mark best/second-best values. FIXED OBJECTS reports box/mask (B/M) at the nominal 16.98k-step budget. Giga floating-point operations (GFLOPs) use one 640^2 image. The speed-retention index (SRI), $F_i G_i/(F_0 G_0)$, uses FPS F and GFLOPs G relative to the No Pretraining (BASE) configuration.

Method	Efficiency			POLYSEG mask		SURFACEMASK mask		FIXED OBJECTS B/M	
	FPS	GFLOPs	SRI	mAP_{50}	$mAP_{50:95}$	mAP_{50}	$mAP_{50:95}$	mAP_{50}	$mAP_{50:95}$
YOLOv9-E [47]	19.6	124	0.58	<u>46.56</u>	<u>27.62</u>	**75.59**	59.68	<u>73.75 / 60.89</u>	<u>53.37</u> / 41.28
YOLOv11-X [16]	**36.1**	160	1.38	46.01	27.41	74.64	59.22	73.00 / 60.14	53.03 / 40.94
OneFormer-Swin-L [12]	7.4	980	1.74	45.72	26.91	75.06	<u>59.91</u>	73.07 / 60.21	53.13 / 41.04
No Pretraining (BASE)	25.8	162	1.00	43.92	26.04	74.73	59.61	72.00 / 59.14	52.54 / 40.45
L2-SP [23]	25.8	164	1.03	44.86	26.61	74.68	59.72	73.08 / 60.41	53.18 / <u>41.52</u>
Full Fine-tuning	25.8	162	1.00	44.25	26.22	74.07	59.31	71.84 / 58.98	52.48 / 40.39
P5 Frozen, plain neck	25.9	162	1.01	42.95	25.32	74.14	59.58	71.22 / 58.36	52.17 / 40.08
RF PRIOR (w/ full decoder)	25.2	393	2.34	**48.42**	**28.25**	<u>75.15</u>	**60.15**	**75.62 / 63.58**	**55.21 / 43.71**

21.81%/16.76% for DFL and 15.94%/9.50% for mask loss; classification improves most proportionally (Figure 5(c)). On FIXED OBJECTS, mask-loss reductions remain positive in both phases, DFL validation reductions remain negative, and the box advantage changes sign (Figure 5(b)). Mean validation DFL changes by +0.00285 for full minus w/o AA alongside lower mask loss, indicating different loss trends across heads that share the adapted feature pyramid.

Cross-target adaptation dynamics. We transfer a YOLOv12-X sidewalk checkpoint's EMA (exponential moving average) encoder after one recorded FIXED OBJECTS adaptation epoch, with fresh necks and heads. Its training excluded the original FIXED OBJECTS validation pool, split here into 768 adaptation/256 evaluation images; Penn-Fudan and WGISD use 120/50 and 110/27 [39, 48]. These short-budget studies use NVIDIA A100 GPUs, 480 AdamW updates [31], three seeds, and original-resolution mask AP; the supplement reports endpoints and hyperpa-

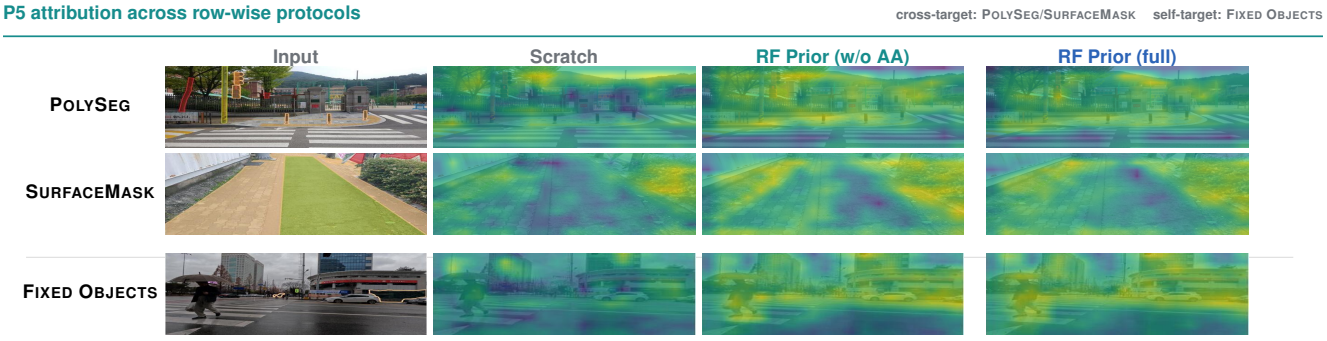


Figure 4. **Post-adaptation P5 gradient-weighted class activation mapping (Grad-CAM)** [40]. Prior-based variants follow *BBox-DS* pretraining and FIXED OBJECTS adaptation; scratch uses FIXED OBJECTS alone. POLYSEG/SURFACEMASK probe retained source-conditioned structure across targets, while FIXED OBJECTS probes self-target object focus. The retained structure motivates Figure 2(c); each row compares attribution on one spatially matched input across models trained under the corresponding adaptation protocol.

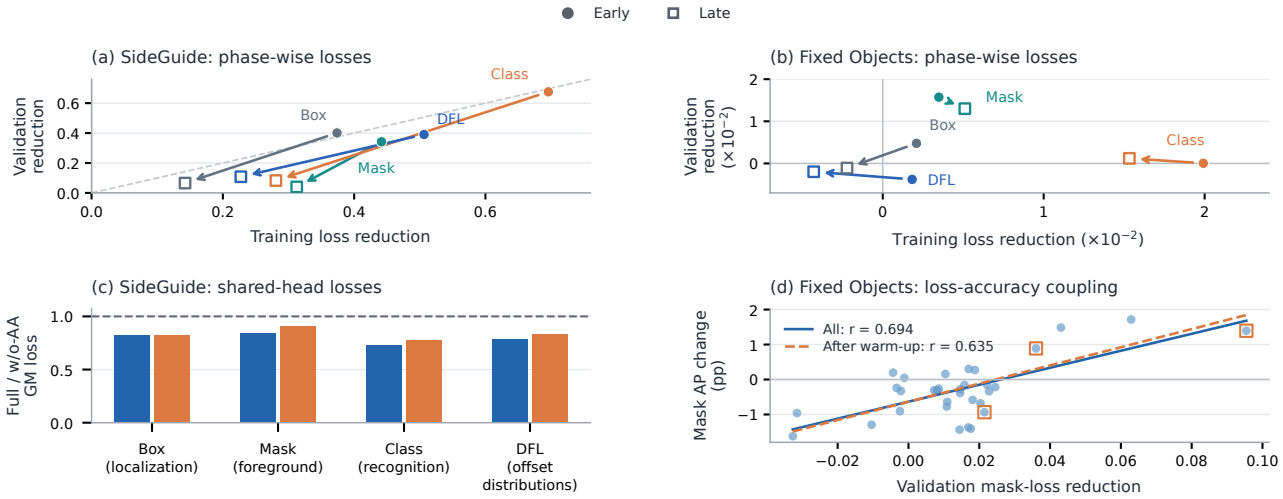


Figure 5. **Complementary optimization diagnostics.** (a,b) Mean train/validation loss reductions (w/o AA minus full), from early circles to late squares: SideGuide splits at 149.8k/299.6k steps; FIXED OBJECTS at 27.17k/56.03k mini-batches. (c) SideGuide geometric-mean (GM) loss ratios: blue training, orange validation. (d) Epoch-matched FIXED OBJECTS mask-loss reduction versus mask AP difference in percentage points (pp); squares mark warm-up epochs 1–3 and lines summarize the observed association across 33 paired epochs.

Table 3. **SideGuide validation.** P/R denote precision/recall and B/M box/mask; subscripts mark second-best checkpoints. Gray marks proposed or extended schedules; within each schedule, bold/underline mark the strongest/second-strongest recorded validation result.

Method	P(B)	R(B)	mAP_{50}(B)	$mAP_{50:95}$(B)	P(M)	R(M)	mAP_{50}(M)	$mAP_{50:95}$(M)
YOLOv9-E [47]	**71.09**$_{70.20}$	49.85$_{50.09}$	55.17$_{54.97}$	41.91$_{41.72}$	**69.69**$_{69.49}$	46.97$_{46.64}$	51.15$_{51.30}$	33.25$_{33.15}$
YOLOv11-X [16]	66.08$_{68.82}$	47.89$_{47.68}$	53.93$_{53.47}$	40.97$_{40.58}$	67.64$_{66.95}$	45.39$_{45.59}$	50.58$_{50.15}$	33.03$_{32.83}$
No Pretraining	66.61$_{66.28}$	47.15$_{46.65}$	52.30$_{52.04}$	39.51$_{39.31}$	65.41$_{64.72}$	44.70$_{44.64}$	49.08$_{48.86}$	32.08$_{31.95}$
⌞ No Pretraining, 50 epochs	67.95$_{67.62}$	49.61$_{49.95}$	54.72$_{54.57}$	41.84$_{41.71}$	69.54$_{69.27}$	45.66$_{45.83}$	51.25$_{51.12}$	33.53$_{33.47}$
Full Fine-tuning	69.38$_{68.91}$	44.97$_{44.88}$	51.53$_{51.15}$	38.78$_{38.42}$	67.72$_{66.85}$	43.11$_{43.01}$	48.30$_{47.92}$	31.48$_{31.22}$
P5 Frozen, plain neck	63.58$_{64.10}$	46.77$_{46.09}$	51.61$_{51.26}$	38.99$_{38.71}$	63.38$_{66.56}$	44.04$_{43.48}$	48.41$_{48.13}$	31.57$_{31.36}$
RF PRIOR (w/ full decoder)	70.22$_{69.48}$	**51.55**$_{51.49}$	**56.95**$_{56.63}$	**42.09**$_{41.84}$	68.08$_{67.95}$	**48.91**$_{48.08}$	**52.86**$_{52.56}$	**33.81**$_{33.71}$
⌞ **RF PRIOR (w/ full decoder), 50 epochs**	**75.84**$_{74.52}$	**56.69**$_{57.08}$	**63.61**$_{63.21}$	**48.26**$_{47.96}$	**75.11**$_{73.96}$	**52.99**$_{52.87}$	**58.98**$_{58.66}$	**37.87**$_{37.64}$

rameters. Full FT, P5 freezing, L2-SP, and DoRA use the native neck, while both RF variants freeze P5 and use SPPF; w/o AA substitutes a local C2f block for attention. Table 5 couples stage selection with representation continuity: P5-frozen variants yield feature distances of 0.663–0.684, versus 0.750/0.751 for native P3/P4 freezing and 0.758 after P5 release. Among the P4-trainable freezing controls, full RF Prior has the smallest P4 displacement (0.151).

Table 4. **FIXED OBJECTS archived checkpoints.** P/R: precision/recall; B/M: box/mask. Subscripts: second-best checkpoints; gray: RF PRIOR variants. Full/w/o-AA use epochs 11/31 and 10/30 under the nominal budgets. Figure 5(d) compares matched epochs.

Budget	Method	P(B)	R(B)	mAP_{50}(B)	$mAP_{50:95}$(B)	P(M)	R(M)	mAP_{50}(M)	$mAP_{50:95}$(M)
16.98k	No Pretraining	$79.77_{76.98}$	$64.28_{62.06}$	$72.00_{68.26}$	$52.54_{49.32}$	$69.54_{66.48}$	$55.01_{54.47}$	$59.14_{57.76}$	$40.45_{38.38}$
16.98k	L2-SP [23]	$79.91_{77.43}$	$65.31_{63.52}$	$73.08_{70.14}$	$53.18_{50.21}$	$69.62_{66.91}$	$55.62_{54.93}$	$60.41_{58.86}$	$41.52_{39.74}$
16.98k	RF PRIOR (w/o AA)	$79.38_{78.14}$	$\mathbf{66.72}_{64.60}$	$73.44_{71.72}$	$53.57_{50.83}$	$\mathbf{73.46}_{70.32}$	$55.78_{55.60}$	$62.27_{60.51}$	$42.98_{40.48}$
16.98k	**RF PRIOR + SPPF + AA**	$\mathbf{80.26}_{75.58}$	$66.50_{67.02}$	$\mathbf{75.62}_{73.07}$	$\mathbf{55.21}_{53.16}$	$69.70_{66.90}$	$\mathbf{57.23}_{55.56}$	$\mathbf{63.58}_{60.45}$	$\mathbf{43.71}_{41.36}$
50.94k	No Pretraining	$\mathbf{87.41}_{86.49}$	$73.83_{74.14}$	$83.36_{83.57}$	$67.01_{66.70}$	$\mathbf{76.53}_{77.66}$	$63.64_{64.19}$	$70.30_{71.33}$	$51.23_{51.48}$
50.94k	L2-SP [23]	$85.96_{85.74}$	$76.92_{76.81}$	$84.72_{84.31}$	$67.89_{67.54}$	$74.91_{75.28}$	$65.91_{66.10}$	$71.42_{71.80}$	$52.21_{52.08}$
50.94k	RF PRIOR (w/o AA)	$83.72_{81.40}$	$\mathbf{81.29}_{81.60}$	$\mathbf{86.64}_{85.98}$	$69.55_{68.67}$	$73.54_{70.91}$	$\mathbf{70.84}_{70.92}$	$\mathbf{73.51}_{73.15}$	$\mathbf{53.86}_{53.77}$
50.94k	**RF PRIOR + SPPF + AA**	$83.99_{84.24}$	$80.13_{79.89}$	$86.48_{86.08}$	$\mathbf{69.84}_{69.38}$	$72.96_{72.81}$	$69.58_{69.19}$	$73.05_{72.64}$	$53.68_{53.47}$

Table 5. **Stage-policy and regularization controls on FIXED OBJECTS.** Three-seed means at update 480; subscripts give standard deviation (SD). $\Delta_3^\theta, \Delta_4^\theta$: relative P3/P4 weight displacement; D_5^{feat}: P5 cosine distance from initialization on 16 fixed images. †: zero imposed by freezing. Native-decoder P5 freeze–release trains P5 during updates 241–480 after the frozen phase.

Model / policy	Δ_3^θ	Δ_4^θ	D_5^{feat}
Native decoder			
Full FT	$0.237_{\pm 0.011}$	$0.212_{\pm 0.003}$	$0.745_{\pm 0.014}$
L2-SP [23]	$0.087_{\pm 0.003}$	$0.009_{\pm 0.001}$	$0.778_{\pm 0.025}$
P3 frozen	$0^\dagger$	$0.207_{\pm 0.015}$	$0.750_{\pm 0.010}$
P4 frozen	$0.233_{\pm 0.008}$	$0^\dagger$	$0.751_{\pm 0.011}$
P5 frozen	$0.227_{\pm 0.011}$	$0.191_{\pm 0.008}$	$0.670_{\pm 0.015}$
P5 freeze–release	$0.232_{\pm 0.015}$	$0.184_{\pm 0.013}$	$0.758_{\pm 0.040}$
RF decoder; P5 frozen			
RF Prior (w/o AA)	$0.242_{\pm 0.008}$	$0.186_{\pm 0.017}$	$0.663_{\pm 0.024}$
RF Prior (full)	$0.243_{\pm 0.003}$	$0.151_{\pm 0.013}$	$0.684_{\pm 0.001}$

L2-SP reduces that displacement to 0.009, while P5 feature distance reaches 0.778, compared with 0.684 for full RF Prior. The regularizer anchors weights, while feature direction depends on upstream adaptation and normalization state, distinguishing parameter proximity from continuity of the target-image representations.

Observation 1 (Readout-dependent P4 adaptation)**.** Across the three targets and three seeds in the author-checkpoint study, the full RF decoder has lower final P4 displacement and late P4 gradient RMS than w/o AA, while mean P3 displacement differs by less than 0.4% on each target.

Final P4 displacements are 0.151/0.147/0.101 for FIXED OBJECTS/Penn-Fudan/WGISD, versus 0.186/0.248/0.130 without AA (Figure 1(b)). Late P4 gradient RMS (updates 241–480) falls by 12.6%/25.1%/16.1% over the matched schedules. In the six-method comparison of Figure 6, full RF Prior keeps P5 parameters fixed while feature cosine distance reaches 0.684/0.724/0.715, distinguishing parameter preservation from adaptive activations. P3 displacement is 0.243/0.246/0.194; its late gradient RMS changes by $+23.8/-13.2/+8.9\%$, contrasting with the consistent P4 reductions. On WGISD, the full decoder has the lowest mean mask loss (2.986) and DFL (1.923), improving on Full FT by 1.7% and 1.4% over the entire schedule (Figure 7). Native Full FT leads the box and classification means, whereas native P5 freezing leads final mask AP at 39.00; full RF Prior gains 0.84 AP over w/o AA. Across targets, full-minus-w/o-AA AP is $+0.35/-3.61/+0.84$; this separates the repeatable change in adaptation dynamics from the target-dependent accuracy ordering at the shared 480-update budget used for all transfer methods.

5. Discussion

Preserved parameters, adaptive representations. Under Definition 1, Proposition 1 links fixed P5 parameters to adaptive upstream features through the contextual readout. Figure 4's cross-target structure is consistent with retained source-conditioned organization after the FIXED OBJECTS label shift, while self-target attribution concentrates on labeled infrastructure. Figure 3 complements this view with direct-target spatial evidence across the same hierarchy.

Table 5 further separates weight anchoring from feature continuity: among adaptive-P4 methods, L2-SP restricts P4 movement most, whereas P5-frozen readouts yield smaller P5 feature distances. Definition 2 distinguishes gradient scale, weight displacement, and feature direction; Figure 6 also resolves rescaling. Native P5 freezing shares the parameter constraint but changes the readout; P3/P4 freezing and midpoint P5 release instead alter which stages adapt.

Where contextual readout changes adaptation. Observation 1 places the repeatable AA signature at P4, with lower displacement and late gradients alongside adaptive P3 features. This stage contrast localizes the contextual readout's measured effect within the adapting hierarchy. At the P5-to-P4 interface, AA makes the combination of pooled context and adaptive features input-dependent. Its effect therefore appears in the derivatives seen by P4 as well as in the features presented to the prediction heads. The observed P4 reductions, together with comparatively similar P3 movement, are consistent with a scale-specific redistribution of adaptation through the preserved contex-

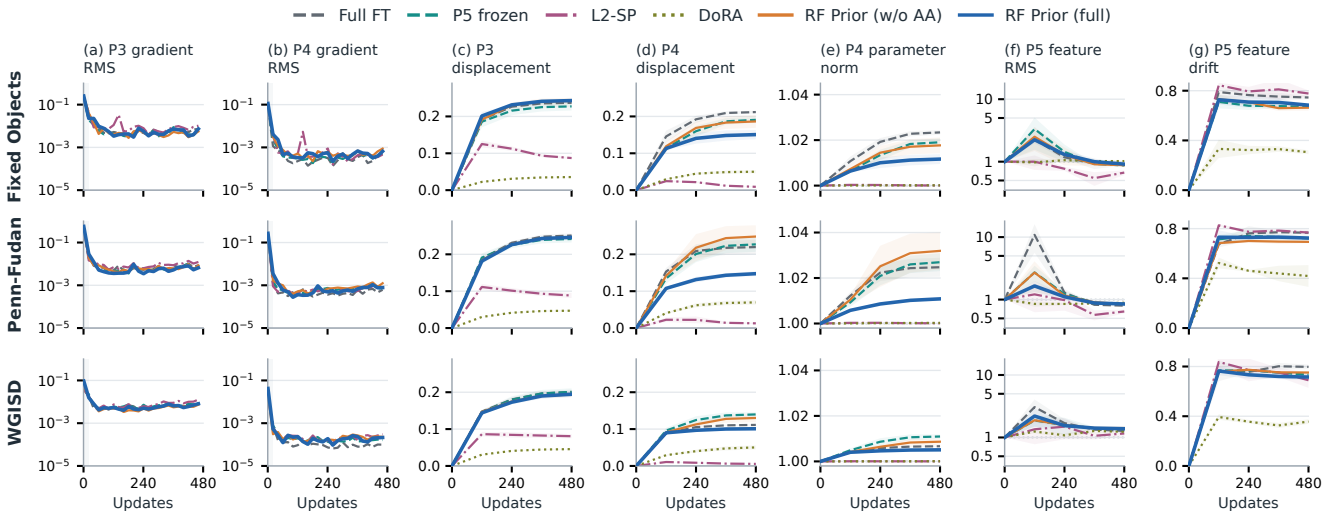


Figure 6. **Gradient, weight, and feature dynamics.** Rows are targets; columns are (a,b) pre-clipping P3/P4 gradient RMS, (c,d) P3/P4 relative displacement, (e) P4 parameter norm and (f) P5 feature RMS relative to initialization, and (g) P5 cosine distance on 16 fixed images. Lines average three seeds over optimizer updates; (c–g) ribbons show SD. Warm-up is shaded; the supplement reports DoRA adapter gradients and complete norm/feature diagnostics across all three pyramid stages over the matched optimization schedules.

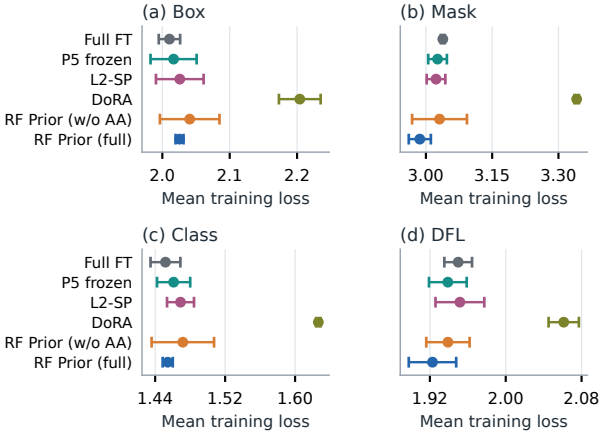


Figure 7. **WGISD shared-head fitting.** Full-schedule means for all four training losses over 480 updates. Points and bars show mean±SD across three seeds, with identical update budgets and shared target annotations for the six evaluated transfer methods.

tual transformation and its trainable readout. At the task level, the PolySeg gain over full fine-tuning (4.17/2.03 mask mAP points) links the design to mask transfer; SurfaceMask complements this result with strict localization under surface-centric labels.

SideGuide and WGISD extend the analysis to distributional box-offset fitting [24], consistent with the adapted pyramid feeding several prediction heads. Mask and box objectives shape P5 inputs through a shared Jacobian, coupling contextual reuse to mask fitting and distributional localization under the same preserved parameter block.

Transfer across supervision and initialization. Holding supervision fixed also permits policy comparisons across target geometries [21, 53]. YOLOv12-S/YOLOv8-S encoders [15, 41] pretrained on Common Objects in Context (COCO) [26] exhibit the same P4 direction across six backbone–target means: AA reduces displacement by 0.9–7.5% and late gradient RMS by 4.5–34.7%, with P3 displacement changes below 2.8%. This repeated P4 pattern connects contextual readout to adaptation across different source initializations and model capacities.

6. Conclusion

RF Prior formulates dense transfer through scale-selective parameter preservation and a trainable contextual readout. Its gradient decomposition explains how preserved P5 parameters participate in target learning as upstream features evolve. This scale-selective design improves strict localization over full fine-tuning across three sidewalk targets at 25.2 FPS. Stage-wise diagnostics identify its characteristic behavior: attention reduces P4 movement and late gradients while P3 remains adaptive and P5 features evolve. Moreover, shared-head and cross-initialization analyses connect this adaptation pattern to different target geometries and source models. The resulting design principle couples the preserved scale to its readout, connecting constrained parameters to adaptive target representations through the gradient structure of instance segmentation when both the target label inventory and instance geometry change.

LLM Usage Statement

OpenAI ChatGPT and Codex were used for manuscript editing, mathematical exposition, literature search, experiment and plotting code, execution support, and project-page preparation. AI visual review assessed the retained validation annotations, as documented in Table 1(b) and the supplementary material. The author performed the final review and validation and takes responsibility for the submitted work.

References

[1] Daniel Bolya, Chong Zhou, Fanyi Xiao, and Yong Jae Lee. YOLACT: Real-time instance segmentation. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 9157–9166, 2019. 2, 3

[2] Nicolas Carion, Francisco Massa, Gabriel Synnaeve, Nicolas Usunier, Alexander Kirillov, and Sergey Zagoruyko. End-to-end object detection with transformers. In *European conference on computer vision*, pages 213–229. Springer, 2020. 2

[3] Bowen Cheng, Ishan Misra, Alexander G Schwing, Alexander Kirillov, and Rohit Girdhar. Masked-attention mask transformer for universal image segmentation. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 1290–1299, 2022. 2

[4] Marius Cordts, Mohamed Omran, Sebastian Ramos, Timo Rehfeld, Markus Enzweiler, Rodrigo Benenson, Uwe Franke, Stefan Roth, and Bernt Schiele. The cityscapes dataset for semantic urban scene understanding. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 3213–3223, 2016. 2

[5] Jifeng Dai, Kaiming He, and Jian Sun. Boxsup: Exploiting bounding boxes to supervise convolutional networks for semantic segmentation. In *Proceedings of the IEEE international conference on computer vision*, pages 1635–1643, 2015. 2

[6] David H Douglas and Thomas K Peucker. Algorithms for the reduction of the number of points required to represent a digitized line or its caricature. *Cartographica: the international journal for geographic information and geovisualization*, 10 (2):112–122, 1973. 2

[7] Yunhui Guo, Honghui Shi, Abhishek Kumar, Kristen Grauman, Tajana Rosing, and Rogerio Feris. Spottune: Transfer learning through adaptive fine-tuning. In *CVPR*, pages 4805–4814, 2019. 1

[8] Kaiming He, Georgia Gkioxari, Piotr Dollár, and Ross Girshick. Mask r-cnn. In *ICCV*, pages 2961–2969, 2017. 2

[9] Lukas Hoyer, Dengxin Dai, and Luc Van Gool. DAFormer: Improving network architectures and training strategies for domain-adaptive semantic segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 9924–9935, 2022. 2

[10] Zilong Huang, Xinggang Wang, Lichao Huang, Chang Huang, Yunchao Wei, and Wenyu Liu. Ccnet: Criss-cross attention for semantic segmentation. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 603–612, 2019. 2

[11] Xinyue Huo, Lingxi Xie, Hengtong Hu, Wengang Zhou, Houqiang Li, and Qi Tian. Domain-agnostic prior for transfer semantic segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 7075–7085, 2022. 2

[12] Jitesh Jain, Jiachen Li, Mang Tik Chiu, Ali Hassani, Nikita Orlov, and Humphrey Shi. Oneformer: One transformer to rule universal image segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 2989–2998, 2023. 2, 4, 5

[13] Menglin Jia, Luming Tang, Bor-Chun Chen, Claire Cardie, Serge Belongie, Bharath Hariharan, and Ser-Nam Lim. Visual prompt tuning. In *European conference on computer vision*, pages 709–727. Springer, 2022. 2

[14] Glenn Jocher. Ultralytics YOLOv5. Software, 2020. 3, 4

[15] Glenn Jocher. Ultralytics YOLOv8: Open-source object detection. GitHub repository: `https://github.com/ultralytics/ultralytics`, 2023. 8

[16] Rahima Khanam and Muhammad Hussain. YOLOv11: An overview of the key architectural enhancements. *arXiv preprint arXiv:2410.17725*, 2024. 5, 6

[17] Anna Khoreva, Rodrigo Benenson, Jan Hosang, Matthias Hein, and Bernt Schiele. Simple does it: Weakly supervised instance and semantic segmentation. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 876–885, 2017. 2

[18] Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete Xiao, Spencer Whitehead, Alexander C Berg, Wan-Yen Lo, et al. Segment anything. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 4015–4026, 2023. 2, 3

[19] James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, et al. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 2017. 1

[20] Shiyi Lan, Zhiding Yu, Christopher Choy, Subhashree Radhakrishnan, Guilin Liu, Yuke Zhu, Larry S. Davis, and Anima Anandkumar. DiscoBox: Weakly supervised instance segmentation and semantic correspondence from box supervision. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 3406–3416, 2021. 2

[21] Yoonho Lee, Annie S. Chen, Fahim Tajwar, Ananya Kumar, Huaxiu Yao, Percy Liang, and Chelsea Finn. Surgical fine-tuning improves adaptation to distribution shifts. In *International Conference on Learning Representations*, 2023. 1, 8

[22] Feng Li, Hao Zhang, Huaizhe Xu, Shilong Liu, Lei Zhang, Lionel M Ni, and Heung-Yeung Shum. Mask dino: Towards a unified transformer-based framework for object detection and segmentation. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 3041–3050, 2023. 2

[23] Xuhong Li, Yves Grandvalet, and Franck Davoine. Explicit inductive bias for transfer learning with convolutional networks. In *Proceedings of the 35th International Conference on Machine Learning*, pages 2825–2834. PMLR, 2018. 1, 2, 4, 5, 7

[24] Xiang Li, Wenhai Wang, Lijun Wu, Shuo Chen, Xiaolin Hu, Jun Li, Jinhui Tang, and Jian Yang. Generalized focal loss: Learning qualified and distributed bounding boxes for dense object detection. In *Advances in Neural Information Processing Systems*, pages 21002–21012, 2020. 3, 4, 8

[25] Zhizhong Li and Derek Hoiem. Learning without forgetting. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 614–629, 2016. 1

[26] Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C. Lawrence Zitnick. Microsoft COCO: Common objects in context. In *European Conference on Computer Vision*, pages 740–755, 2014. 8

[27] Tsung-Yi Lin, Piotr Dollár, Ross Girshick, Kaiming He, Bharath Hariharan, and Serge Belongie. Feature pyramid networks for object detection. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 2117–2125, 2017. 1, 2, 3

[28] Shih-Yang Liu, Chien-Yi Wang, Hongxu Yin, Pavlo Molchanov, Yu-Chiang Frank Wang, Kwang-Ting Cheng, and Min-Hung Chen. DoRA: Weight-decomposed low-rank adaptation. In *Proceedings of the 41st International Conference on Machine Learning*, pages 32100–32121, 2024. 2

[29] Ze Liu, Yutong Lin, Yue Cao, Han Hu, Yixuan Wei, Zheng Zhang, Stephen Lin, and Baining Guo. Swin transformer: Hierarchical vision transformer using shifted windows. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 10012–10022, 2021. 2

[30] Jonathan Long, Evan Shelhamer, and Trevor Darrell. Fully convolutional networks for semantic segmentation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 3431–3440, 2015. 2

[31] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019. 2, 5

[32] Wenjie Luo, Yujia Li, Raquel Urtasun, and Richard Zemel. Understanding the effective receptive field in deep convolutional neural networks. In *NeurIPS*, pages 4898–4906, 2016. 1, 2

[33] Gerhard Neuhold, Tobias Ollmann, Samuel Rota Bulo, and Peter Kontschieder. The mapillary vistas dataset for semantic understanding of street scenes. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 4990–4999, 2017. 2

[34] Kibaek Park, Youngtaek Oh, Soomin Ham, Kyungdon Joo, Hyokyoung Kim, Hyoyoung Kum, and In So Kweon. SideGuide: A large-scale sidewalk dataset for guiding impaired people. In *2020 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 10022–10029. IEEE, 2020. 1, 2, 5

[35] Yansong Peng, Hebei Li, Peixi Wu, Yueyi Zhang, Xiaoyan Sun, and Feng Wu. D-FINE: Redefine regression task of DETRs as fine-grained distribution refinement. In *International Conference on Learning Representations*, 2025. 4

[36] Urs Ramer. An iterative procedure for the polygonal approximation of plane curves. *Computer graphics and image processing*, 1(3):244–256, 1972. 2

[37] Sylvestre-Alvise Rebuffi, Hakan Bilen, and Andrea Vedaldi. Learning multiple visual domains with residual adapters. In *Advances in Neural Information Processing Systems*, 2017. 2

[38] Joseph Redmon, Santosh Divvala, Ross Girshick, and Ali Farhadi. You only look once: Unified, real-time object detection. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 779–788, 2016. 2

[39] Thiago T. Santos, Leonardo L. de Souza, Andreza A. dos Santos, and Sandra Avila. Grape detection, segmentation and tracking using deep neural networks and three-dimensional association. *Computers and Electronics in Agriculture*, 170: 105247, 2020. 2, 5

[40] Ramprasaath R Selvaraju, Michael Cogswell, Abhishek Das, Ramakrishna Vedantam, Devi Parikh, and Dhruv Batra. Grad-cam: Visual explanations from deep networks via gradient-based localization. In *Proceedings of the IEEE international conference on computer vision*, pages 618–626, 2017. 6

[41] Yunjie Tian, Qixiang Ye, and David Doermann. YOLOv12: Attention-centric real-time object detectors. In *Advances in Neural Information Processing Systems*, 2025. 1, 3, 4, 8

[42] Zhi Tian, Chunhua Shen, Xinlong Wang, and Hao Chen. Boxinst: High-performance instance segmentation with box annotations. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 5443–5452, 2021. 2

[43] Cristina Vasconcelos, Vighnesh Birodkar, and Vincent Dumoulin. Proper reuse of image classification features improves object detection. In *CVPR*, pages 13628–13637, 2022. 1, 2

[44] Tuan-Hung Vu, Himalaya Jain, Maxime Bucher, Matthieu Cord, and Patrick Perez. ADVENT: Adversarial entropy minimization for domain adaptation in semantic segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 2517–2526, 2019. 2

[45] Ao Wang, Hui Chen, Lihao Liu, Kai Chen, Zijia Lin, Jungong Han, and Guiguang Ding. YOLOv10: Real-time end-to-end object detection. *Advances in Neural Information Processing Systems*, 37:107984–108011, 2024. 2

[46] Chien-Yao Wang, Alexey Bochkovskiy, and Hong-Yuan Mark Liao. YOLOv7: Trainable bag-of-freebies sets new state-of-the-art for real-time object detectors. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 7464–7475, 2023.

[47] Chien-Yao Wang, I-Hau Yeh, and Hong-Yuan Mark Liao. YOLOv9: Learning what you want to learn using programmable gradient information. In *European Conference on Computer Vision*, pages 1–21, 2024. 2, 5, 6

[48] Liming Wang, Jianbo Shi, Gang Song, and I-fan Shen. Object detection combining recognition and segmentation. In *Asian Conference on Computer Vision*, pages 189–199, 2007. 2, 5

[49] Xiaolong Wang, Ross Girshick, Abhinav Gupta, and Kaiming He. Non-local neural networks. In *Proceedings of the*

IEEE conference on computer vision and pattern recognition, pages 7794–7803, 2018. 2

[50] Xinlong Wang, Rufeng Zhang, Tao Kong, Lei Li, and Chunhua Shen. Solov2: Dynamic and fast instance segmentation. In *NeurIPS*, pages 17721–17732, 2020. 2

[51] Zhixiang Wei, Lin Chen, Yi Jin, Xiaoxiao Ma, Tianle Liu, Pengyang Ling, Ben Wang, Huaian Chen, and Jinjin Zheng. Stronger, fewer, & superior: Harnessing vision foundation models for domain generalized semantic segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 28619–28630, 2024. 2

[52] Enze Xie, Wenhai Wang, Zhiding Yu, Anima Anandkumar, Jose M Alvarez, and Ping Luo. Segformer: Simple and efficient design for semantic segmentation with transformers. *Advances in neural information processing systems*, 34: 12077–12090, 2021. 2

[53] Jason Yosinski, Jeff Clune, Yoshua Bengio, and Hod Lipson. How transferable are features in deep neural networks? *Advances in neural information processing systems*, 27, 2014. 1, 2, 8

[54] Fisher Yu, Haofeng Chen, Xin Wang, Wenqi Xian, Yingying Chen, Fangchen Liu, Vashisht Madhavan, and Trevor Darrell. BDD100K: A diverse driving dataset for heterogeneous multitask learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 2636–2645, 2020. 2

[55] Chaoning Zhang, Dongshen Han, Yu Qiao, Jung Uk Kim, Sung-Ho Bae, Seungkyu Lee, and Choong Seon Hong. Faster segment anything: Towards lightweight SAM for mobile applications. *arXiv preprint arXiv:2306.14289*, 2023. 4

[56] H. Zhao, J. Shi, X. Qi, X. Wang, and J. Jia. Pyramid scene parsing network. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 2881–2890, 2017. 2

[57] Yian Zhao, Wenyu Lv, Shangliang Xu, Jinman Wei, Guanzhong Wang, Qingqing Dang, Yi Liu, and Jie Chen. DETRs beat YOLOs on real-time object detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 16965–16974, 2024. 2

[58] Sixiao Zheng, Jiachen Lu, Hengshuang Zhao, Xiatian Zhu, Zekun Luo, Yabiao Wang, Yanwei Fu, Jianfeng Feng, Tao Xiang, Philip HS Torr, et al. Rethinking semantic segmentation from a sequence-to-sequence perspective with transformers. In *CVPR*, pages 6881–6890, 2021. 2

[59] Zhaohui Zheng, Ping Wang, Wei Liu, Jinze Li, Rongguang Ye, and Dongwei Ren. Distance-IoU loss: Faster and better learning for bounding box regression. In *AAAI*, pages 12993–13000, 2020. 3